

Kaggle Titanic

Titanic Survivor Prediction

Data loading and organization

```
library(caret)

## Loading required package: lattice
## Loading required package: ggplot2
## Registered S3 methods overwritten by 'ggplot2':
##   method      from
##   [.quosures   rlang
##   c.quosures   rlang
##   print.quosures rlang

library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(MASS)

##
## Attaching package: 'MASS'
## The following object is masked from 'package:dplyr':
##
##   select

train <- read.csv("Datasets/train.csv", stringsAsFactors = TRUE, na.strings = "")
test <- read.csv("Datasets/test.csv", stringsAsFactors = TRUE)

test$Survived <- NA

dat <- rbind(train,test)

summary(dat)

##   PassengerId      Survived      Pclass
##   Min.   :    1   Min.   :0.0000   Min.   :1.000
##   1st Qu.:  328   1st Qu.:0.0000   1st Qu.:2.000
##   Median :  655   Median :0.0000   Median :3.000
##   Mean   :  655   Mean   :0.3838   Mean   :2.295
##   3rd Qu.:  982   3rd Qu.:1.0000   3rd Qu.:3.000
##   Max.   : 1309   Max.   :1.0000   Max.   :3.000
```

```
##          NA's      :418
##          Name          Sex          Age
## Connolly, Miss. Kate      : 2   female:466   Min.      : 0.17
## Kelly, Mr. James          : 2   male  :843   1st Qu.:21.00
## Abbing, Mr. Anthony       : 1                      Median :28.00
## Abbott, Mr. Rossmore Edward : 1                      Mean   :29.88
## Abbott, Mrs. Stanton (Rosa Hunt): 1                 3rd Qu.:39.00
## Abelson, Mr. Samuel       : 1                      Max.   :80.00
## (Other)                   :1301                   NA's   :263
##      SibSp      Parch      Ticket      Fare
## Min.      :0.0000   Min.      :0.000   CA. 2343: 11   Min.      : 0.000
## 1st Qu.:0.0000   1st Qu.:0.000   1601      : 8   1st Qu.: 7.896
## Median :0.0000   Median :0.000   CA 2144   : 8   Median : 14.454
## Mean   :0.4989   Mean   :0.385   3101295   : 7   Mean   : 33.295
## 3rd Qu.:1.0000   3rd Qu.:0.000   347077    : 7   3rd Qu.: 31.275
## Max.   :8.0000   Max.   :9.000   347082    : 7   Max.   :512.329
##                                     (Other) :1261   NA's     :1
##      Cabin      Embarked
##      :327      C      :270
## C23 C25 C27      : 6      Q      :123
## B57 B59 B63 B66: 5      S      :914
## G6              : 5      NA's: 2
## B96 B98          : 4
## (Other)          :275
## NA's             :687
```

```
#convert survived and pclass to factor variable
dat$Survived <- as.factor(dat$Survived)
dat$Pclass <- as.factor(dat$Pclass)

#Cabin Na values -> 0, otherwise 1
dat$Cabin <- as.factor(ifelse(is.na(dat$Cabin), 0, 1))

#Gender -> male = 0, female = 1
dat$Sex <- as.factor(ifelse(dat$Sex == "male", 0, 1))

#NA values -> mean
dat$Age[is.na(dat$Age)] <- mean(dat$Age, na.rm=TRUE)
dat$Fare[is.na(dat$Fare)] <- mean(dat$Fare, na.rm=TRUE)
dat$Embarked[is.na(dat$Embarked)] <- "S"

#family size (if family = 1, then it's alone)
dat$family <- dat$SibSp + dat$Parch + 1

#converting names
name <- dat$Name
name <- sub(".*, ", "", name)
name <- sub("\\\\..*", "", name)
head(name)
```

```
## [1] "Mr" "Mrs" "Miss" "Mrs" "Mr" "Mr"
```

```

dat$Name <- as.character(name)

dat$Name[dat$Name %in% c("Capt","Col","Don","Dr","Jonkheer", "Lady", "Major","Mlle","Mme","Ms","Rev","S")]

dat$Name <- as.factor(dat$Name)
table(dat$Name)

##
## Master    Miss      Mr      Mrs unknown
##      61     260     757     197      34

#dropping variables
dat <- subset(dat, select = -c(PassengerId, SibSp, Parch, Ticket))

names(dat)

## [1] "Survived" "Pclass"  "Name"      "Sex"       "Age"       "Fare"
## [7] "Cabin"    "Embarked" "family"

summary(dat)

## Survived  Pclass      Name      Sex      Age
## 0 :549    1:323  Master : 61    0:843  Min.   : 0.17
## 1 :342    2:277  Miss  :260    1:466  1st Qu.:22.00
## NA's:418  3:709  Mr     :757      Median :29.88
##                               Mrs     :197      Mean    :29.88
##                               unknown: 34      3rd Qu.:35.00
##                               Max.     :80.00
##
##      Fare      Cabin  Embarked  family
## Min.   : 0.000    0:687    C:270   Min.   : 1.000
## 1st Qu.: 7.896    1:622    Q:123   1st Qu.: 1.000
## Median :14.454      S:916   Median : 1.000
## Mean   :33.295      Mean   : 1.884
## 3rd Qu.:31.275      3rd Qu.: 2.000
## Max.   :512.329      Max.   :11.000

train1 <- dat[1:891,] #train.csv
test1 <- dat[892:1309,] #test.csv

set.seed(125)

split<-createDataPartition(y=train1$Survived,p=0.7,list=FALSE)

training <- train1[split,] #training set of train data
testing <- train1[-split,] #test set of train data, which is validation set

#test1 is the test dataset

```

Data Exploration

```

library(dplyr)
library(purrr)

##

```

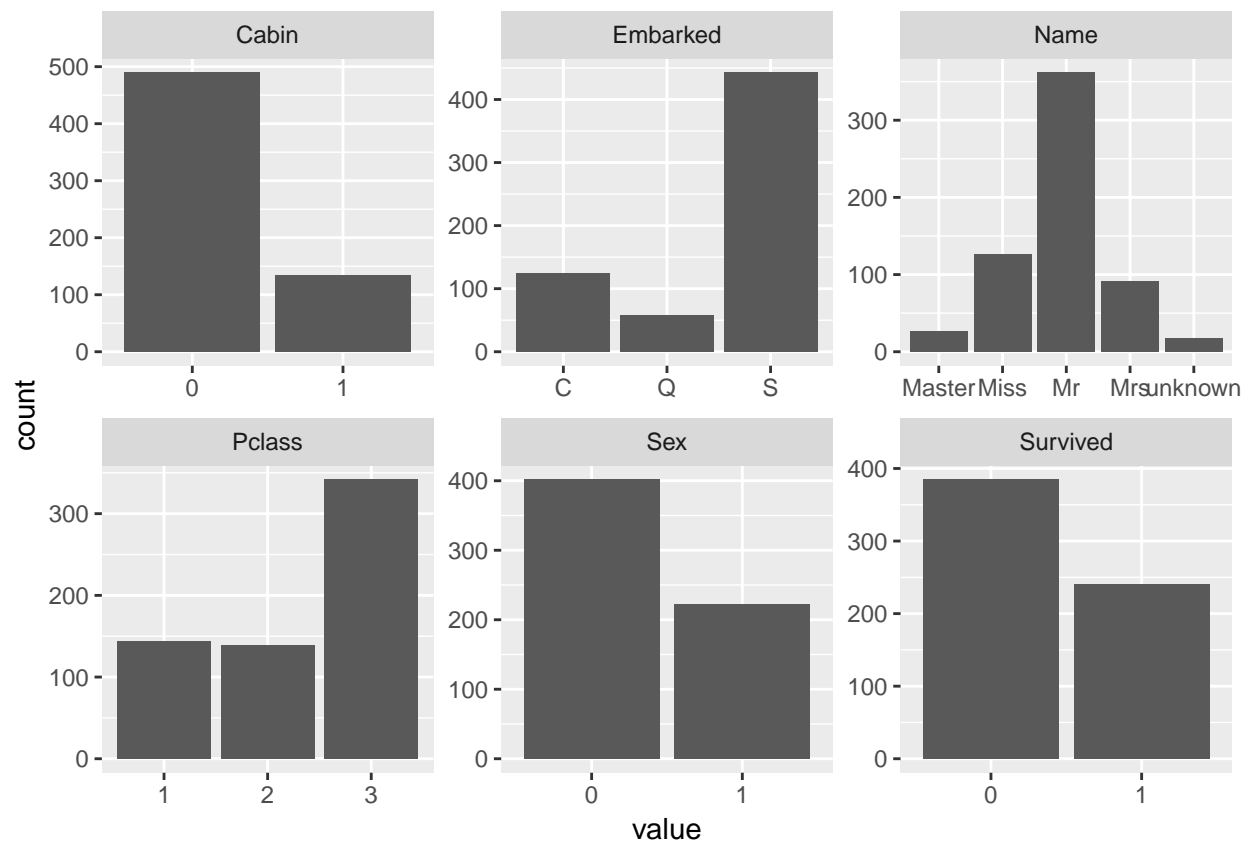
```
## Attaching package: 'purrr'

## The following object is masked from 'package:caret':
##
## lift
```


```
library(tidyr)
library(ggplot2)
```

```
###all variables
training %>% keep(is.factor) %>% gather() %>%
  ggplot(aes(value)) +
  facet_wrap(~key, scales="free")+
  geom_bar()
```

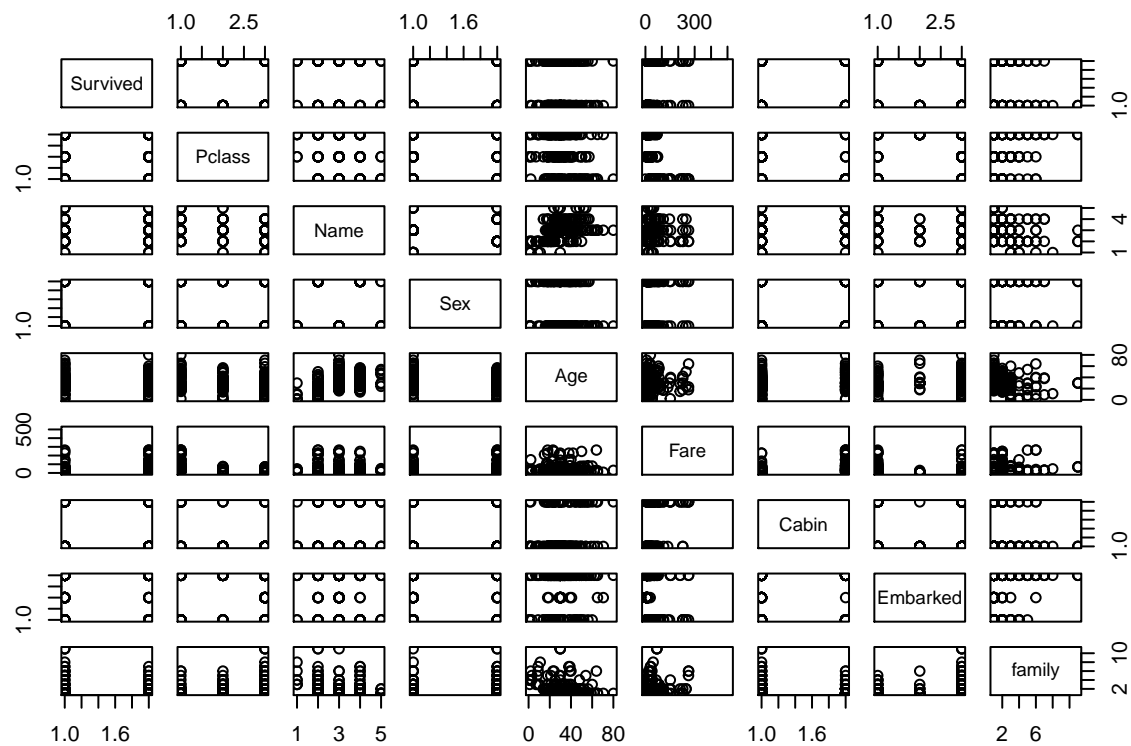
```
## Warning: attributes are not identical across measure variables;
## they will be dropped
```



```
training %>% keep(is.numeric) %>% gather() %>%
  ggplot(aes(value)) +
  facet_wrap(~key, scales="free")+
  geom_histogram(bins=30)+
  geom_density()
```



```
pairs(training,col=train$Survived)
```



Logistic Regression

```
library(boot)

##
## Attaching package: 'boot'
## The following object is masked from 'package:lattice':
##
##      melanoma

library(ROCR)

## Loading required package: gplots
##
## Attaching package: 'gplots'
## The following object is masked from 'package:stats':
##
##      lowess

logistic.fit <- glm(Survived ~ ., data=training, family=binomial)

summary(logistic.fit)

##
## Call:
## glm(formula = Survived ~ ., family = binomial, data = training)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3109  -0.5738  -0.4029   0.5231   2.4922
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   4.052945   0.891576   4.546 5.47e-06 ***
## Pclass2      -0.283878   0.498261  -0.570 0.568856
## Pclass3      -1.295398   0.499055  -2.596 0.009440 **
## NameMiss     -16.581174  639.499709  -0.026 0.979314
## NameMr        -3.749260   0.651305  -5.757 8.59e-09 ***
## NameMrs      -15.786362  639.499776  -0.025 0.980306
## Nameunknown  -3.735382   0.973263  -3.838 0.000124 ***
## Sex1         15.702310  639.499405   0.025 0.980411
## Age          -0.019111   0.010537  -1.814 0.069724 .
## Fare          0.003516   0.003241   1.085 0.278094
## Cabin1        0.860558   0.425625   2.022 0.043190 *
## EmbarkedQ     -0.072022   0.462585  -0.156 0.876273
## EmbarkedS     -0.392620   0.299850  -1.309 0.190403
## family        -0.502213   0.107321  -4.680 2.88e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 832.49  on 624  degrees of freedom
## Residual deviance: 512.66  on 611  degrees of freedom
```

```
## AIC: 540.66
##
## Number of Fisher Scoring iterations: 14
logistic.probs <- predict(logistic.fit, newdata = testing, type="response")

logistic.pred <- ifelse(logistic.probs > 0.5, 1, 0)

test.pred <- testing$Survived

table(logistic.pred, test.pred)

##           test.pred
## logistic.pred  0   1
##           0 148  26
##           1  16  76


sum(diag(table(logistic.pred, test.pred)))/nrow(testing)

## [1] 0.8421053
#accuracy = 84.21%

#cv error with 10 folds
cv.glm(training, logistic.fit, K=10)$delta[1]

## [1] 0.1336628
#around 13% error rate

ROCpred <- prediction(logistic.pred, test.pred)
ROCperf <- performance(ROCpred, "tpr", "fpr")
plot(ROCperf, colorize = TRUE, print.cutoffs.at = seq(0.1,0.1), text.adj = c(-0.2,1.7))
```



```

performance(ROCpred, "auc")@y.values

## [[1]]
## [1] 0.8237685
#AUC value = 0.8237

confusionMatrix(as.factor(logistic.pred), test.pred)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##      0 148   26
##      1   16   76
##
##              Accuracy : 0.8421
##              95% CI : (0.7926, 0.8838)
##      No Information Rate : 0.6165
##      P-Value [Acc > NIR] : 6.804e-16
##
##              Kappa : 0.6598
##
##      McNemar's Test P-Value : 0.1649
##
##              Sensitivity : 0.9024
##              Specificity : 0.7451
##      Pos Pred Value : 0.8506
##      Neg Pred Value : 0.8261
##      Prevalence : 0.6165
##      Detection Rate : 0.5564
##      Detection Prevalence : 0.6541
##      Balanced Accuracy : 0.8238
##
##      'Positive' Class : 0
##

```

Logistic Regression Accuracy -> 84.21%


Decision Tree

```

library(tree)

tree.training <- tree(Survived~., training)
plot(tree.training);text(tree.training,pretty=0)

```

```
summary(tree.training)
```

```
##
## Classification tree:
## tree(formula = Survived ~ ., data = training)
## Variables actually used in tree construction:
## [1] "Name" "Fare" "Pclass" "family" "Cabin"
## Number of terminal nodes: 8
## Residual mean deviance: 0.7791 = 480.7 / 617
## Misclassification error rate: 0.1632 = 102 / 625
```

```
tree.pred=predict(tree.training,testing,type="class") #with test datasets
```

```
table(tree.pred , testing$Survived)
```

```
##
## tree.pred  0  1
##           0 147 28
##           1  17 74
```

```
sum(diag(table(tree.pred, testing$Survived)))/nrow(testing)
```

```
## [1] 0.8308271
```

```
#accuracy 83.08%
```

```
deviance(tree.training)
```

```
## [1] 480.6808
```

```
misclass.tree(tree.training)
```

```
## [1] 102
```


```
#pruning via cv
```

```
cv.training=cv.tree(tree.training,FUN=prune.misclass)
cv.training
```


```
## $size
## [1] 8 6 4 2 1
```

```
##
## $dev
## [1] 110 110 114 151 240
##
## $k
## [1] -Inf 0.0 4.0 13.5 103.0
##
## $method
## [1] "misclass"
##
## attr("class")
## [1] "prune" "tree.sequence"
```


```
plot(cv.training)
```



```
plot(cv.training$size, cv.training$dev, type="b")
```



```
prune.training <- prune.misclass(tree.training,best=7)#actually same with tree
plot(prune.training);text(prune.training,pretty=0)
```



```
#pruned prediction
tree.pred <- predict(prune.training, testing, type="class")
table(tree.pred, testing$Survived)
```

```
##
## tree.pred  0  1
##           0 147 28
##           1  17 74
```

```
sum(diag(table(tree.pred, testing$Survived)))/nrow(testing)
```

```
## [1] 0.8308271
```

```
#83.08%
```

```
class(tree.pred)
```

```
## [1] "factor"
```


```
class(test.pred)
```

```
## [1] "factor"
```

```
ROCpred <- prediction(as.numeric(tree.pred), test.pred)
```

```
ROCperf <- performance(ROCpred, "tpr", "fpr")
```

```
plot(ROCperf, colorize = TRUE, print.cutoffs.at = seq(0.1,0.1), text.adj = c(-0.2,1.7))
```



```
performance(ROCpred, "auc")@y.values
```

```
## [[1]]
```

```
## [1] 0.8109158
```

```
#auc = 0.8109
```

```
confusionMatrix(tree.pred, test.pred)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  0    1
```

```
##           0 147  28
```

```
##           1   17  74
```

```
##
```

```
##           Accuracy : 0.8308
```

```
##           95% CI : (0.7803, 0.8738)
```

```
## No Information Rate : 0.6165
```

```
## P-Value [Acc > NIR] : 2.211e-14
```

```
##
```

```
##                Kappa : 0.6348
##
## Mcnemar's Test P-Value : 0.136
##
##          Sensitivity : 0.8963
##          Specificity : 0.7255
##          Pos Pred Value : 0.8400
##          Neg Pred Value : 0.8132
##          Prevalence : 0.6165
##          Detection Rate : 0.5526
##          Detection Prevalence : 0.6579
##          Balanced Accuracy : 0.8109
##
##          'Positive' Class : 0
##
```

Decision Tree Accuracy -> 83.08%

Random Forest

```
library(randomForest)
```

```
## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:dplyr':
##
##      combine
## The following object is masked from 'package:ggplot2':
##
##      margin
```

```
library(MASS)
```

```
#base random forest
rf.training <- randomForest(Survived~., data=training, importance=TRUE)
rf.training

##
## Call:
## randomForest(formula = Survived ~ ., data = training, importance = TRUE)
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 2
##
##              OOB estimate of  error rate: 17.44%
## Confusion matrix:
##      0    1 class.error
## 0 348  37  0.0961039
```

```
## 1 72 168 0.3000000
```

```
#2 mtry - classification defaults = sqrt(p) = sqrt(8) ~= 2, n.tree = 500, oob error rate = around 17%
```

```
#find best mtry
```

```
oob.err=double(8)
```

```
test.err=double(8)
```

```
for(mtry in 1:8){
```

```
  fit <- randomForest(Survived~., data=training, mtry=mtry, ntree=500)
```

```
  oob.err[mtry] <- 1-sum(diag(table(training$Survived,fit$predicted)))/nrow(training)
```

```
  pred <- predict(fit, testing)
```

```
  test.err[mtry] <- 1-sum(diag(table(testing$Survived,pred)))/nrow(testing)
```


```
  cat(mtry," ")
```

```
}
```

```
## 1 2 3 4 5 6 7 8
```

```
matplot(1:mtry,cbind(test.err,oob.err),pch=19,col=c("red","blue"),type="b",ylab="Error Rate")
```

```
legend("topright",legend=c("Test","OOB"),pch=19,col=c("red","blue"))
```



```
#2 mtry might be the best
```

```
#updated random forest
```

```
rf.training.update <- randomForest(Survived~., data=training, mtry=2, importance=TRUE, ntree=500)
```

```
rf.training.update
```

```
##
```

```
## Call:
```

```
## randomForest(formula = Survived ~ ., data = training, mtry = 2, importance = TRUE, ntree = 500)
```

```
## Type of random forest: classification
```

```
## Number of trees: 500
```

```
## No. of variables tried at each split: 2
```

```
##
```

```

##          OOB estimate of  error rate: 17.44%
## Confusion matrix:
##      0   1 class.error
## 0 349  36  0.09350649
## 1   73 167  0.30416667

rf.pred <- predict(rf.training.update, testing)
table(testing$Survived, rf.pred)

##      rf.pred
##           0    1
##    0 153   11
##    1   30   72

sum(diag(table(testing$Survived, rf.pred)))/nrow(testing)

## [1] 0.8458647

#around 84.58% accuracy
#since random forest provides different accuracy in every try, I made a function to get a mean value of

mean.acc <- function(training, testing, mtry, ntree, try, formula){
  acc <- NULL

  for(i in 1:try){
    randomFtrain <- randomForest(formula, data=training, mtry=mtry, importance=TRUE, ntree=ntree)
    randomFpred <- predict(randomFtrain, testing)
    acc[i] <- sum(diag(table(testing$Survived, randomFpred)))/nrow(testing)
  }

  return(mean(acc))
}

#2 mtry, 500 ntree, 10 tries
mean.acc(training = training, testing = testing, mtry = 2, ntree = 500, try = 10, formula = as.formula(

## [1] 0.8454887

#around 84.5% accuracy


importance(rf.training.update)

##           0           1 MeanDecreaseAccuracy MeanDecreaseGini
## Pclass    18.134499  20.011172             27.53240           19.545816
## Name       23.064960  20.999153             27.16987           47.525815
## Sex        23.792787  20.190242             26.82020           38.822613
## Age        11.570976   7.040736             14.19494           28.820272
## Fare       14.556035  15.535657             23.20695           37.421976
## Cabin      12.280560   1.776714             13.72102           11.226495
## Embarked    6.638619   7.246604             10.31469            6.311426
## family     21.015530   5.218679             20.89440           18.407316

varImpPlot(rf.training.update)

```

rf.training.update



#we can improve our RF model by removing Embarked

```
rf.update <- randomForest(Survived~.-Embarked, data=training, mtry=2, importance=TRUE, ntree=500)
rf.update
```

```
##
## Call:
## randomForest(formula = Survived ~ . - Embarked, data = training,          mtry = 2, importance = TRUE, ntree = 500)
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 2
##
## OOB estimate of  error rate: 16.8%
## Confusion matrix:
##      0   1 class.error
## 0 353  32  0.08311688
## 1   73 167  0.30416667
```

```
rf.pred <- predict(rf.update, testing)
table(testing$Survived, rf.pred)
```

```
##      rf.pred
##         0    1
## 0 152  12
## 1   30  72
```

```
sum(diag(table(testing$Survived,rf.pred)))/nrow(testing)
```

```
## [1] 0.8421053
```

#around 84%


```
mean.acc(training, testing, mtry = 2, ntree = 500, try = 10, formula = as.formula(Survived~.-Embarked))
```

```
## [1] 0.8421053
```


```
#accruacy around 84.3%
```

```
importance(rf.update)
```

```
##           0           1 MeanDecreaseAccuracy MeanDecreaseGini
## Pclass 19.281870 20.462255          27.74701          20.47984
## Name   26.192644 23.155114          29.69098          52.75031
## Sex    22.439350 19.452844          24.80557          35.32776
## Age     7.367302  8.111400          11.06260          27.34985
## Fare    14.755373 17.864728          23.97459          38.27450
## Cabin   12.514360  1.968030          13.22990          10.53989
## family  24.767579  7.058626          25.66010          19.07239
```

```
varImpPlot(rf.update)
```

rf.update




```
class(rf.pred)
```

```
## [1] "factor"
```

```
ROCpred <- prediction(as.numeric(rf.pred), test.pred)
```

```
ROCperf <- performance(ROCpred, "tpr", "fpr")
```

```
plot(ROCperf, colorize = TRUE, print.cutoffs.at = seq(0.1,0.1), text.adj = c(-0.2,1.7))
```



```
performance(ROCpred, "auc")@y.values
```

```
## [[1]]
## [1] 0.8163558
```

```
#auc ~= around 0.8212
```

```
confusionMatrix(rf.pred, test.pred)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 152  30
##           1  12  72
##
##               Accuracy : 0.8421
##               95% CI : (0.7926, 0.8838)
##           No Information Rate : 0.6165
##           P-Value [Acc > NIR] : 6.804e-16
##
##               Kappa : 0.6545
##
##  Mcnemar's Test P-Value : 0.008712
##
##           Sensitivity : 0.9268
##           Specificity : 0.7059
##           Pos Pred Value : 0.8352
##           Neg Pred Value : 0.8571
##           Prevalence : 0.6165
##           Detection Rate : 0.5714
##           Detection Prevalence : 0.6842
##           Balanced Accuracy : 0.8164
```

```
##
##      'Positive' Class : 0
##
```

Random Forest Accuracy -> around 84.5%


boosting

```
library(gbm)
```

```
## Loaded gbm 2.1.5
```

```
set.seed(123)
boost.training <- gbm(as.character(Survived)~., data=training,
                      distribution = "bernoulli", n.trees = 500,
                      interaction.depth = 3, shrinkage = 0.01, cv.folds=5)

summary(boost.training)
```



```
##      var    rel.inf
## Name      Name 45.234382
## Fare      Fare 14.554858
## Pclass    Pclass 13.760218
## family    family 9.136101
## Age       Age  5.878103
## Sex       Sex  4.864853
## Cabin     Cabin 4.525450
## Embarked  Embarked 2.046035
```

```
#test on the training set
boost.pred <- predict(boost.training, newdata = training, n.trees = 500, type = "response") #prob
boost.pred <- ifelse(boost.pred > 0.5, 1, 0)
```

```

table(training$Survived, boost.pred)

##      boost.pred
##      0      1
## 0 356   29
## 1   54  186

sum(diag(table(training$Survived, boost.pred)))/nrow(training)

## [1] 0.8672
#86.08% accuracy

#test on the test set
boost.pred <- predict(boost.training, newdata = testing, n.trees = 500, type = "response")

boost.pred <- ifelse(boost.pred > 0.5, 1, 0)

table(testing$Survived, boost.pred)

##      boost.pred
##      0      1
## 0 148   16
## 1   29   73

sum(diag(table(testing$Survived, boost.pred)))/nrow(testing)

## [1] 0.8308271
#83.83% accuracy

#cross validation to get the best ntree
folds <- sample(rep(1:5, length = nrow(training)))
folds

## [1] 4 2 1 3 3 2 4 2 4 3 3 1 2 1 2 3 2 5 1 3 2 5 1 3 1 3 1 5 4 2 2 3 5 1 4
## [36] 5 2 2 5 2 1 4 1 3 4 3 1 3 5 4 2 5 2 3 1 1 1 1 3 5 1 4 5 4 1 4 1 4 1 1
## [71] 5 5 3 2 2 2 3 1 2 5 2 1 2 1 4 5 2 5 4 5 1 4 2 5 3 4 3 1 5 1 2 3 1 4 1
## [106] 3 2 5 3 3 2 5 4 5 5 5 5 5 3 3 2 3 2 3 2 4 5 1 4 1 5 1 5 4 3 5 1 3 3 5
## [141] 4 5 2 5 3 3 1 4 1 1 2 2 4 5 1 2 3 1 3 2 4 5 1 3 4 5 5 5 1 3 3 1 4 5 1
## [176] 4 1 4 1 5 4 2 1 5 1 2 3 3 2 5 4 2 2 2 4 1 3 5 3 5 2 1 5 1 1 4 4 5 2 4
## [211] 4 4 4 2 1 4 2 5 4 4 4 3 1 5 3 1 4 1 1 5 4 2 5 3 3 2 2 5 4 2 3 1 2 5 1
## [246] 2 3 4 3 4 4 5 5 5 2 3 2 4 3 1 5 3 1 4 5 5 2 1 3 2 3 3 3 4 2 5 5 2 4 5
## [281] 5 3 2 2 1 4 5 2 1 4 1 3 5 1 5 1 4 3 3 3 4 3 5 4 4 1 1 1 3 3 2 5 4 4 1
## [316] 5 3 3 4 4 2 4 1 4 3 1 3 3 3 5 4 5 4 3 1 3 4 4 1 4 1 5 4 2 2 5 1 5 4 3
## [351] 3 3 5 2 2 5 3 1 3 5 4 1 1 3 5 3 3 2 2 2 3 5 4 4 2 3 2 5 4 2 3 1 1 2 4
## [386] 2 5 5 3 4 4 4 1 2 5 2 3 5 1 2 1 3 3 1 4 2 5 1 2 3 4 4 3 1 3 2 3 2 3 5
## [421] 1 2 3 4 1 1 1 5 1 3 1 2 5 3 5 3 1 2 2 3 5 2 4 5 2 3 1 3 4 2 3 2 3 1 1
## [456] 5 1 4 4 4 3 4 1 2 2 5 2 1 4 3 1 3 1 5 4 5 5 4 2 2 2 2 4 5 5 5 2 3 5 3
## [491] 3 4 3 1 2 2 4 4 5 4 5 2 1 2 3 4 2 3 1 2 1 2 5 1 3 2 2 1 4 5 3 2 2 4 5
## [526] 2 4 5 2 5 2 1 5 5 1 1 2 4 5 2 4 4 5 1 3 1 4 2 5 1 4 4 3 3 3 4 1 1 1 4
## [561] 5 2 2 1 4 1 1 5 4 3 4 3 4 5 2 3 3 2 4 3 1 1 2 5 2 3 5 5 4 5 4 1 4 5 1
## [596] 5 4 4 4 1 2 1 1 5 2 3 4 3 4 5 2 2 4 2 5 2 5 3 3 5 4 5 2 3 4

```

```

table(folds)

## folds
##   1   2   3   4   5
## 125 125 125 125 125

ntree <- seq(500, 4000, by=500)
cv.errors <- matrix(NA, length(ntree),5)
cv.errors <- cbind(ntree, cv.errors)
cv.errors

##      ntree
## [1,]   500 NA NA NA NA NA
## [2,]  1000 NA NA NA NA NA
## [3,]  1500 NA NA NA NA NA
## [4,]  2000 NA NA NA NA NA
## [5,]  2500 NA NA NA NA NA
## [6,]  3000 NA NA NA NA NA
## [7,]  3500 NA NA NA NA NA
## [8,]  4000 NA NA NA NA NA

for(i in 1:5){
  for(j in ntree){
    boost.fit <- gbm(as.character(Survived)~., data = training[folds != i,], distribution = "bernoulli")
    boost.pred <- predict(boost.fit, training[folds == i,], n.trees = j)
    boost.pred <- ifelse(boost.pred > 0.5, 1, 0)
    cv.errors[which(j==ntree), i+1] = 1- sum(diag(table(training$Survived[folds==i], boost.pred)))/nrow
  }
}

cv.errors

##      ntree
## [1,]   500 0.240 0.144 0.200 0.184 0.152
## [2,]  1000 0.232 0.128 0.208 0.176 0.160
## [3,]  1500 0.248 0.136 0.208 0.160 0.168
## [4,]  2000 0.232 0.136 0.208 0.160 0.168
## [5,]  2500 0.256 0.144 0.216 0.160 0.160
## [6,]  3000 0.248 0.128 0.216 0.168 0.144
## [7,]  3500 0.256 0.136 0.232 0.176 0.152
## [8,]  4000 0.256 0.144 0.224 0.176 0.168

cv.errors1 <- cv.errors[, -1]
cv.errors1


##
## [1,] 0.240 0.144 0.200 0.184 0.152
## [2,] 0.232 0.128 0.208 0.176 0.160
## [3,] 0.248 0.136 0.208 0.160 0.168
## [4,] 0.232 0.136 0.208 0.160 0.168
## [5,] 0.256 0.144 0.216 0.160 0.160
## [6,] 0.248 0.128 0.216 0.168 0.144
## [7,] 0.256 0.136 0.232 0.176 0.152
## [8,] 0.256 0.144 0.224 0.176 0.168

cv.error.rates=sqrt(apply(cv.errors1,1,mean))
which.min(cv.error.rates)

```


```
## [1] 2
```

```
plot(cv.error.rates,pch=19,type="b")
```



```
boost.training <- gbm(as.character(Survived)~., data=training, distribution = "bernoulli", n.trees = 1000)
```

```
summary(boost.training)
```



```
##          var  rel.inf
## Name      Name 37.914412
## Fare      Fare 18.366654
## Pclass    Pclass 12.655526
```

```
## Age          Age 10.205115
## family      family 8.173176
## Sex         Sex 5.696653
## Cabin       Cabin 4.223646
## Embarked    Embarked 2.764818

#test on the test set
boost.pred <- predict(boost.training, newdata = testing, n.trees = 1000, type = "response")
boost.pred <- ifelse(boost.pred > 0.5, 1, 0)

table(testing$Survived, boost.pred)

##      boost.pred
##      0      1
## 0 147   17
## 1   25   77

sum(diag(table(testing$Survived, boost.pred)))/nrow(testing)

## [1] 0.8421053

#84.21% accuracy
```

Boosting -> 84.21%

SVM - linear

```
library(e1071)
set.seed(123)
svm.fit <- svm(Survived~., data=training, scale=FALSE, kernel="linear", cost=5)
print(svm.fit)

##
## Call:
## svm(formula = Survived ~ ., data = training, kernel = "linear",
##      cost = 5, scale = FALSE)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:   5
##   gamma:    0.07142857
##
## Number of Support Vectors: 258

summary(svm.fit)

##
## Call:
## svm(formula = Survived ~ ., data = training, kernel = "linear",
##      cost = 5, scale = FALSE)
##
##
## Parameters:
##   SVM-Type:  C-classification
```

```

## SVM-Kernel: linear
## cost: 5
## gamma: 0.07142857
##
## Number of Support Vectors: 258
##
## ( 130 128 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1
tune.out <- tune(svm, Survived~., data=training, kernel = "linear", ranges = list(cost = c(0.001, 0.01,
summary(tune.out)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
## cost
## 10
##
## - best performance: 0.1646953
##
## - Detailed performance results:
## cost error dispersion
## 1 0.001 0.3808244 0.04926884
## 2 0.010 0.2030978 0.04264397
## 3 0.100 0.1967486 0.04079913
## 4 1.000 0.1710701 0.03838288
## 5 5.000 0.1695341 0.02972506
## 6 10.000 0.1646953 0.03213910
## 7 15.000 0.1646953 0.03213910
## 8 20.000 0.1646953 0.03213910
tune.out$best.model

##
## Call:
## best.tune(method = svm, train.x = Survived ~ ., data = training,
## ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 15, 20)),
## kernel = "linear")
##
##
## Parameters:
## SVM-Type: C-classification
## SVM-Kernel: linear
## cost: 10
## gamma: 0.07142857
##

```



```

## Number of Support Vectors: 255
svm.best <- tune.out$best.model
summary(svm.best)

##
## Call:
## best.tune(method = svm, train.x = Survived ~ ., data = training,
##   ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 15, 20)),
##   kernel = "linear")
##
##
## Parameters:
##   SVM-Type: C-classification
##   SVM-Kernel: linear
##       cost: 10
##   gamma: 0.07142857
##
## Number of Support Vectors: 255
##
## ( 129 126 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1

svm.pred <- predict(svm.best, testing, type="class")
table(predict = svm.pred, truth = testing$Survived)


##          truth
## predict    0    1
##          0 148  28
##          1  16  74

sum(diag(table(svm.pred, testing$Survived)))/nrow(testing)

## [1] 0.8345865
#83.46 accuracy

ROCpred <- prediction(as.numeric(svm.pred), test.pred)
ROCperf <- performance(ROCpred, "tpr", "fpr")
plot(ROCperf, colorize = TRUE, print.cutoffs.at = seq(0.1,0.1), text.adj = c(-0.2,1.7))

```



```
performance(ROCpred, "auc")@y.values
```

```
## [[1]]
## [1] 0.8139646
```

```
#auc = 0.8139
```

```
confusionMatrix(svm.pred, testing$Survived)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  0    1
```

```
##           0 148  28
```

```
##           1  16  74
```

```
##
```

```
##           Accuracy : 0.8346
```

```
##           95% CI : (0.7844, 0.8772)
```

```
## No Information Rate : 0.6165
```

```
## P-Value [Acc > NIR] : 7.118e-15
```

```
##
```

```
##           Kappa : 0.6422
```

```
##
```

```
## McNemar's Test P-Value : 0.09725
```

```
##
```

```
##           Sensitivity : 0.9024
```

```
##           Specificity : 0.7255
```

```
## Pos Pred Value : 0.8409
```

```
## Neg Pred Value : 0.8222
```

```
## Prevalence : 0.6165
```

```
## Detection Rate : 0.5564
```

```
## Detection Prevalence : 0.6617
```

```
## Balanced Accuracy : 0.8140
```

```
##  
##      'Positive' Class : 0  
##
```

SVM linear Accuracy -> 83.46%

SVM - radial

```
svm.fit <- svm(Survived~., data=training, scale=FALSE, kernel="radial", cost=5)  
print(svm.fit)
```

```
##  
## Call:  
## svm(formula = Survived ~ ., data = training, kernel = "radial",  
##      cost = 5, scale = FALSE)  
##  
##  
## Parameters:  
##      SVM-Type:  C-classification  
##      SVM-Kernel: radial  
##           cost:  5  
##          gamma: 0.07142857  
##  
## Number of Support Vectors: 445
```

```
summary(svm.fit)
```

```
##  
## Call:  
## svm(formula = Survived ~ ., data = training, kernel = "radial",  
##      cost = 5, scale = FALSE)  
##  
##  
## Parameters:  
##      SVM-Type:  C-classification  
##      SVM-Kernel: radial  
##           cost:  5  
##          gamma: 0.07142857  
##  
## Number of Support Vectors: 445  
##  
## ( 215 230 )  
##  
##  
## Number of Classes: 2  
##  
## Levels:  
## 0 1
```

```
tune.out <- tune(svm, Survived~., data=training, kernel = "radial", ranges = list(cost = c(0.001, 0.01,
```

```
summary(tune.out)
```

```
##  
## Parameter tuning of 'svm':
```

```

##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.1664363
##
## - Detailed performance results:
##   cost      error dispersion
## 1  0.001 0.3841014 0.06229791
## 2  0.010 0.3841014 0.06229791
## 3  0.100 0.1935996 0.05389965
## 4  1.000 0.1664363 0.04209901
## 5  5.000 0.1711982 0.04505016
## 6 10.000 0.1711214 0.04533815
## 7 15.000 0.1790835 0.05147086
## 8 20.000 0.1742960 0.05290177

```

```

tune.out$best.model

```

```

##
## Call:
## best.tune(method = svm, train.x = Survived ~ ., data = training,
##   ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 15, 20)),
##   kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##     cost:    1
##     gamma:   0.07142857
##
## Number of Support Vectors: 299

```

```

svm.best <- tune.out$best.model
summary(svm.best)

```

```

##
## Call:
## best.tune(method = svm, train.x = Survived ~ ., data = training,
##   ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 15, 20)),
##   kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##     cost:    1
##     gamma:   0.07142857
##
## Number of Support Vectors: 299
##

```

```
## ( 146 153 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1


svm.pred <- predict(svm.best, testing, type="class")
table(predict = svm.pred, truth = testing$Survived)

##          truth
## predict    0    1
##          0 149  28
##          1  15  74

sum(diag(table(svm.pred, testing$Survived)))/nrow(testing)

## [1] 0.8383459
#83.83% accuracy

ROCpred <- prediction(as.numeric(svm.pred), test.pred)
ROCperf <- performance(ROCpred, "tpr", "fpr")
plot(ROCperf, colorize = TRUE, print.cutoffs.at = seq(0.1,0.1), text.adj = c(-0.2,1.7))
```



```
performance(ROCpred, "auc")@y.values

## [[1]]
## [1] 0.8170134
#auc = 0.8170

confusionMatrix(svm.pred, testing$Survived)
```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 149  28
##           1   15  74
##
##           Accuracy : 0.8383
##           95% CI : (0.7885, 0.8805)
##       No Information Rate : 0.6165
##       P-Value [Acc > NIR] : 2.231e-15
##
##           Kappa : 0.6497
##
##  Mcnemar's Test P-Value : 0.06725
##
##           Sensitivity : 0.9085
##           Specificity : 0.7255
##       Pos Pred Value : 0.8418
##       Neg Pred Value : 0.8315
##           Prevalence : 0.6165
##       Detection Rate : 0.5602
##       Detection Prevalence : 0.6654
##       Balanced Accuracy : 0.8170
##
##       'Positive' Class : 0
##
SVM - radial Accuracy -> 83.83%
Logistic Regression Accuracy -> 84.21%
Decision Tree Accuracy -> 83.08%
Random Forest Accuracy -> around 84.5%
Boosting -> 84.21%
SVM linear Accuracy -> 83.46%
SVM radial Accuracy -> 83.83%

```